

# Utkarsh Tripathi

[ ]:

[ ]:

## Credit Risk Modelling

- Credit risk is an important topic in the field of finance because banks and other financial institutions heavily invest in reducing their credit risk.
- There is an element of risk when these financial institutions lend money to clients. So ideally to reduce their risk and loan money only to clients they are sure are going to pay them back.

**Project goal:** The objective of this project is to build a model to predict probability of a client defaulting a loan.

- Defaulting on a loan is the failure of a borrower to pay the principal or interest on a security or loan.
- The dataset has been collected from the UCI machine learning repository
- **The following steps will be followed in building the model.**
- Data preparation and Pre-processing
- Exploratory Data Analysis
- Feature Engineering and Selection
- Model Development and Model Evaluation
- Hyperparameter Tuning

```
[2]: #importing necessary packages
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from matplotlib import rcParams
rcParams['figure.figsize']=(12,5)
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
sns.set_style('whitegrid')
pd.set_option('display.max_columns',None)

from collections import Counter
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import f1_score , confusion_matrix , classification_report,roc_auc_score , roc_curve
from sklearn.model_selection import StratifiedKFold , cross_val_score
from sklearn.model_selection import GridSearchCV
```

```
[3]: #importing the data
data = pd.read_excel('Credit_Risk_PD_Model_Dataset.xlsx')
data.head()
```

	Industry	CASH_AND_CASH_EQUIVALENTS_floored1	COST_OF_GOODS_SOLD_1	CURRENT_ASSETS_1	CURRENT_LIABILITIES_1	DEPRECIATION_1	EBITDA_1	FIXED
0	Wholesale	323773.0		NaN	476381.0	516192.0	78247.0	357842.0
1	EduArtsHealthSocial	18593.0		11753.0	21913.0	257405.0	6988.0	42274.0
2	Retail	202720.0		3420276.0	665190.0	373352.0	76899.0	406223.0
3	Agriculture - Other	87515.0		7905311.0	1820195.0	2153519.0	702651.0	675750.0
4	EduArtsHealthSocial	272938.0		237933.0	381958.0	681622.0	NaN	-515748.0

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 573 entries, 0 to 572
Data columns (total 39 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   Industry         573 non-null    object  
 1   CASH_AND_CASH_EQUIVALENTS_floored1 573 non-null    float64 
 2   COST_OF_GOODS_SOLD_1      518 non-null    float64 
 3   CURRENT_ASSETS_1        573 non-null    float64 
 4   CURRENT_LIABILITIES_1    573 non-null    float64 
 5   DEPRFCATTON_1          534 non-null    float64 
 6   EBITDA_1                573 non-null    float64 
 7   FIXED                   573 non-null    float64 
 8   FTE                     573 non-null    float64 
 9   GROSSPROFIT_1           573 non-null    float64 
 10  INDUSTRY               573 non-null    float64 
 11  INVENTORY              573 non-null    float64 
 12  LIABILITIES            573 non-null    float64 
 13  MARGIN                  573 non-null    float64 
 14  NETINCOME               573 non-null    float64 
 15  OPERATINGCOSTS          573 non-null    float64 
 16  OPERATINGINCOME          573 non-null    float64 
 17  OPERATINGINCOMEPERCENT  573 non-null    float64 
 18  OPERATINGINCOMERATE     573 non-null    float64 
 19  OPERATINGINCOMERATEPERC 573 non-null    float64 
 20  OPERATINGINCOMERATEPERC 573 non-null    float64 
 21  OPERATINGINCOMERATEPERC 573 non-null    float64 
 22  OPERATINGINCOMERATEPERC 573 non-null    float64 
 23  OPERATINGINCOMERATEPERC 573 non-null    float64 
 24  OPERATINGINCOMERATEPERC 573 non-null    float64 
 25  OPERATINGINCOMERATEPERC 573 non-null    float64 
 26  OPERATINGINCOMERATEPERC 573 non-null    float64 
 27  OPERATINGINCOMERATEPERC 573 non-null    float64 
 28  OPERATINGINCOMERATEPERC 573 non-null    float64 
 29  OPERATINGINCOMERATEPERC 573 non-null    float64 
 30  OPERATINGINCOMERATEPERC 573 non-null    float64 
 31  OPERATINGINCOMERATEPERC 573 non-null    float64 
 32  OPERATINGINCOMERATEPERC 573 non-null    float64 
 33  OPERATINGINCOMERATEPERC 573 non-null    float64 
 34  OPERATINGINCOMERATEPERC 573 non-null    float64 
 35  OPERATINGINCOMERATEPERC 573 non-null    float64 
 36  OPERATINGINCOMERATEPERC 573 non-null    float64 
 37  OPERATINGINCOMERATEPERC 573 non-null    float64 
 38  OPERATINGINCOMERATEPERC 573 non-null    float64 
```

```

6   EBITDA_1           573 non-null  float64
7   FIXED_AND_NON_CURRENT_ASSETS_1    561 non-null  float64
8   GROSS_PROFIT_1          573 non-null  float64
9   INTEREST_EXPENSE_1        480 non-null  float64
10  LONG_TERM_DEBT_1          359 non-null  float64
11  NET_FIXED_ASSETS_1         555 non-null  float64
12  NET_INCOME_1             573 non-null  float64
13  NET_INTANGIBLES_1          573 non-null  float64
14  NET_TRADE_RECV_1          475 non-null  float64
15  NON_CURRENT_LIABILITIES_1    449 non-null  float64
16  RETAINED_EARNINGS_1         573 non-null  float64
17  TOTAL_ASSETS_1             573 non-null  float64
18  TOTAL_DEBT_1              527 non-null  float64
19  TOTAL_INVENTORY_1          481 non-null  float64
20  TURNOVER_1                573 non-null  float64
21  WORKING_CAPITAL_1          573 non-null  float64
22  Debt_to_TNW_1              527 non-null  float64
23  Equity_to_liability_1       573 non-null  float64
24  Equity_to_Assets_1          573 non-null  float64
25  Long_Term_Debt_to_Equity_1    359 non-null  float64
26  NetfixedAssets_to_TotalDebt_1 515 non-null  float64
27  Total_Liab_to_Total_Assets_1 573 non-null  float64
28  EBITDA_to_total_Debt_1       527 non-null  float64
29  EBITDA_to_Interest_1         480 non-null  float64
30  EBITDA_to_curr_liab_1        573 non-null  float64
31  Cash_to_Total_assets_1        573 non-null  float64
32  Current_Ratio_1             573 non-null  float64
33  Liquidity_Ratio_1            475 non-null  float64
34  Curr_to_Total_Assets_1        573 non-null  float64
35  Gross_Margin_perc_1          518 non-null  float64
36  Net_Income_Margin_1          573 non-null  float64
37  Total_Liab_to_Total_Assets_2_1 562 non-null  float64
38  Default                      573 non-null  float64
dtypes: float64(38), object(1)
memory usage: 174.7+ KB

```

```
[5]: #checking for the distribution(count) of target classes
data['Default'] = data['Default'].astype('int')
ax = sns.countplot(data['Default'])
plt.title('distribution(count) of target classes')
for p in ax.patches:
    ax.annotate('{:.1f}'.format(p.get_height()), (p.get_x() + 0.25, p.get_height() + 0.01))
```



```
[6]: #checking for missing values in the data
round(data.isna().sum()/len(data)*100,2).sort_values(ascending=False)
```

```

[6]: Long_Term_Debt_to_Equity_1      37.35
LONG_TERM_DEBT_1                   37.35
NON_CURRENT_LIABILITIES_1          21.64
Liquidity_Ratio_1                  17.10
NET_TRADE_RECV_1                   17.10
INTEREST_EXPENSE_1                 16.23
EBITDA_to_Interest_1               16.23
TOTAL_INVENTORY_1                  16.06
NetfixedAssets_to_TotalDebt_1       10.12
Gross_Margin_perc_1                 9.60
COST_OF_GOODS SOLD_1                9.00
TOTAL_DEBT_1                        8.03
EBITDA_to_total_Debt_1              8.03
Debt_to_TNW_1                        8.03
DEPRECIATION_1                      6.81
NET_FIXED_ASSETS_1                  3.14
FIXED_AND_NON_CURRENT_ASSETS_1       2.09
Total_Liab_to_Total_Assets_2_1       1.92
Curr_to_Total_Assets_1                0.00
Current_Ratio_1                      0.00
Net_Income_Margin_1                  0.00
Cash_to_Total_assets_1                0.00
EBITDA_to_curr_liab_1                0.00
Total_Liab_to_Total_Assets_1          0.00
Industry                            0.00
Equity_to_Assets_1                   0.00
Equity_to_liability_1                 0.00
WORKING_CAPITAL_1                     0.00
TURNOVER_1                           0.00
CASH_AND_CASH_EQUIVALENTS_floored1    0.00
TOTAL_ASSETS_1                        0.00
RETAINED_EARNINGS_1                   0.00

```

```

NET_INNANGIBLES_1          0.00
NET_INCOME_1                0.00
GROSS_PROFIT_1                0.00
EBITDA_1                      0.00
CURRENT_LIABILITIES_1        0.00
CURRENT_ASSETS_1                0.00
Default                          0.00
dtype: float64

```

- considering a threshold of 25% for the missing values and,
- dropping the features which are having missing values more than the threshold

## Removing features with more than 25% missing values

```
[7]: #dropping missing values
data.dropna(axis=1, thresh=round(len(data)-len(data)/100*25), inplace=True)
```

```
[8]: #checking for the distribution of independent variables
data.describe().apply(lambda s : s.apply('{0:.5f}'.format))
```

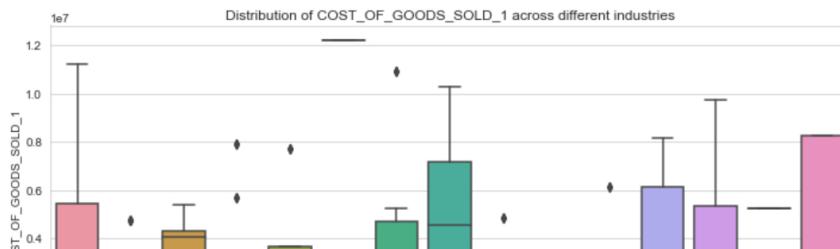
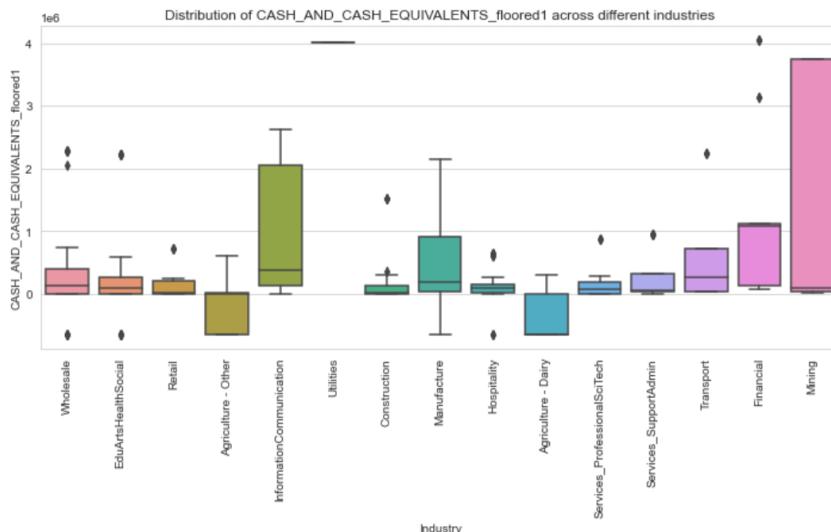
	CASH_AND_CASH_EQUIVALENTS_floored1	COST_OF_GOODS_SOLD_1	CURRENT_ASSETS_1	CURRENT LIABILITYES_1	DEPRECIATION_1	EBITDA_1	FIXED_AND_NON
count	573.00000	518.00000	573.00000	573.00000	534.00000	573.00000	
mean	260737.46422	2418543.23359	1882030.51832	1519872.27051	95440.10300	354311.38220	
std	821575.22736	2862968.43545	3882481.64782	3559609.21764	155056.63976	852420.84859	
min	-644364.00000	5116.00000	1936.00000	228.00000	271.00000	-7324453.00000	
25%	790.00000	315268.00000	236199.00000	253171.00000	14941.00000	73034.00000	
50%	49605.00000	1024016.00000	770200.00000	706407.00000	45219.00000	189612.00000	
75%	284872.00000	4245791.75000	2180921.00000	1656906.00000	111614.25000	489874.00000	
max	4052236.00000	12212492.00000	34522271.00000	32716574.00000	1334359.00000	7660463.00000	

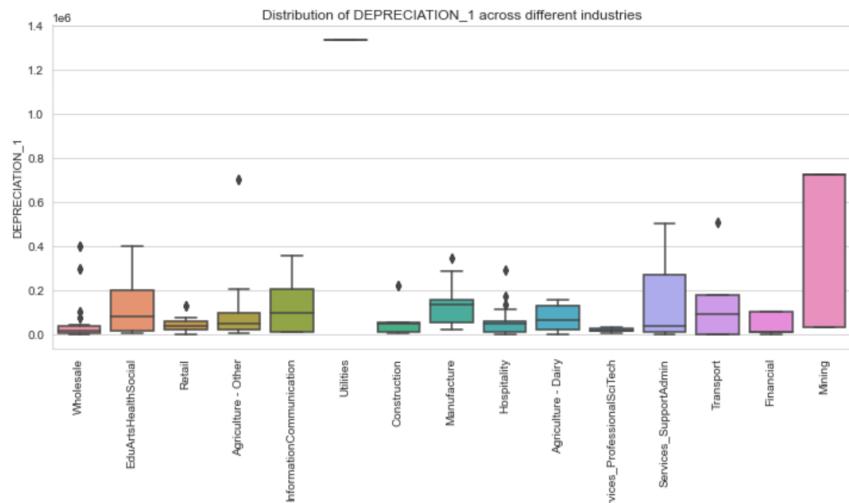
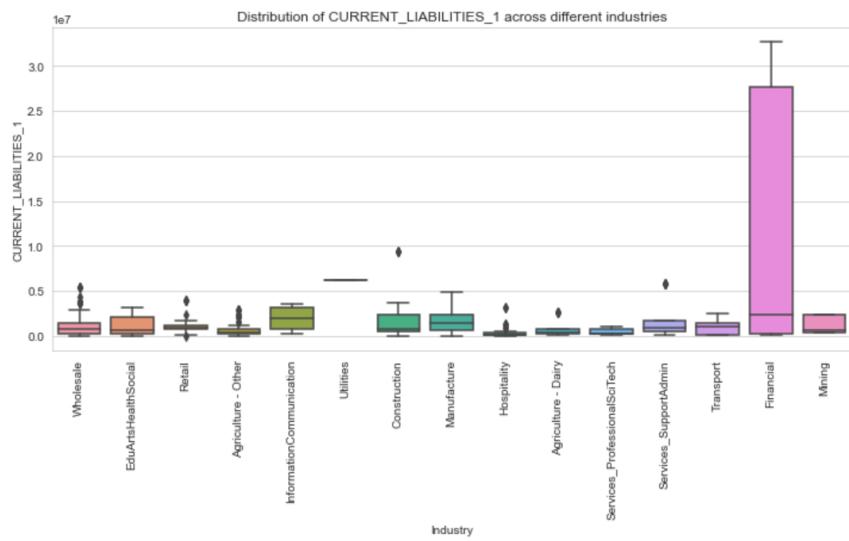
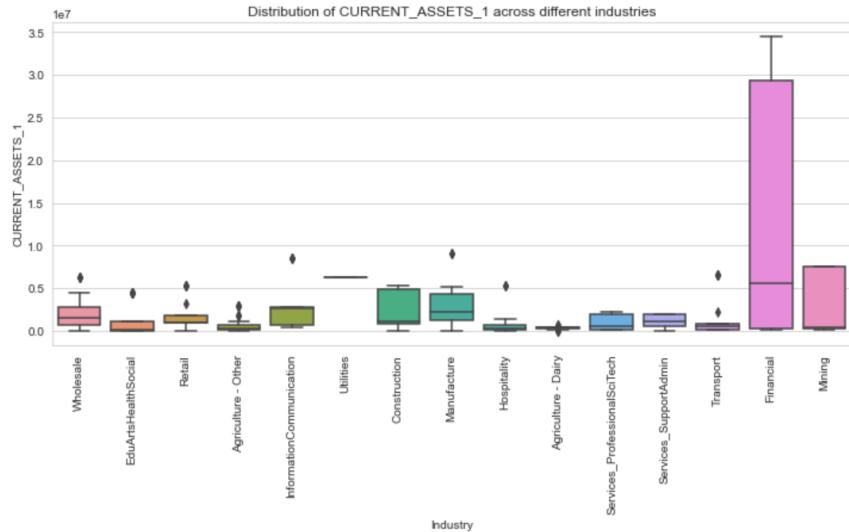
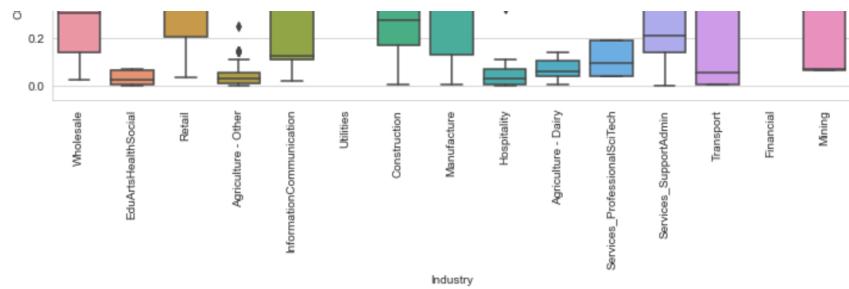
## Checking for the distribution of the features with respect to the different industries

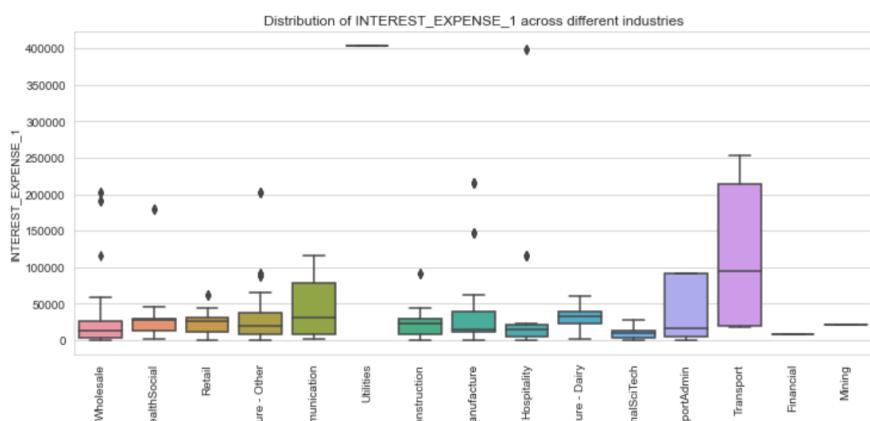
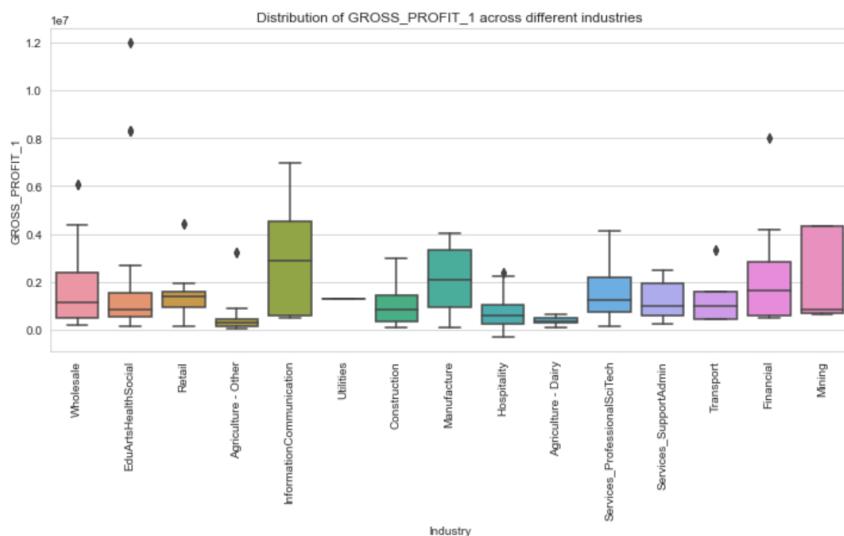
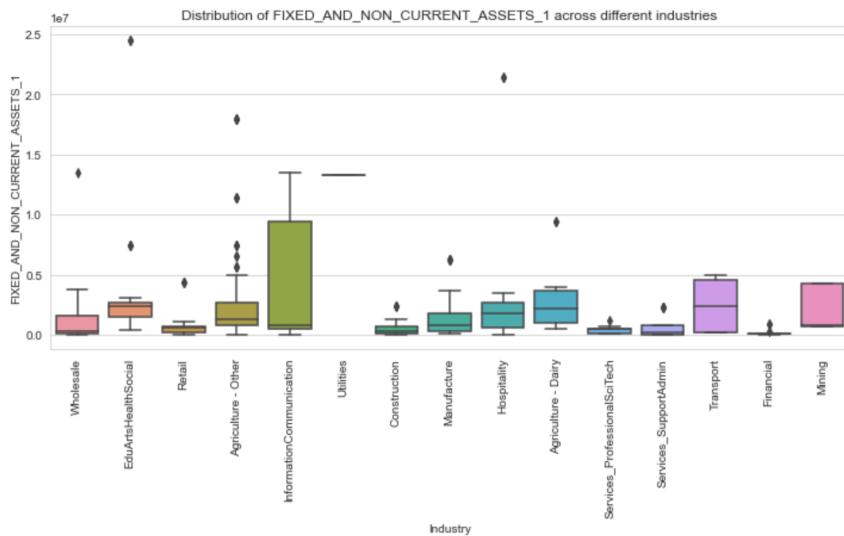
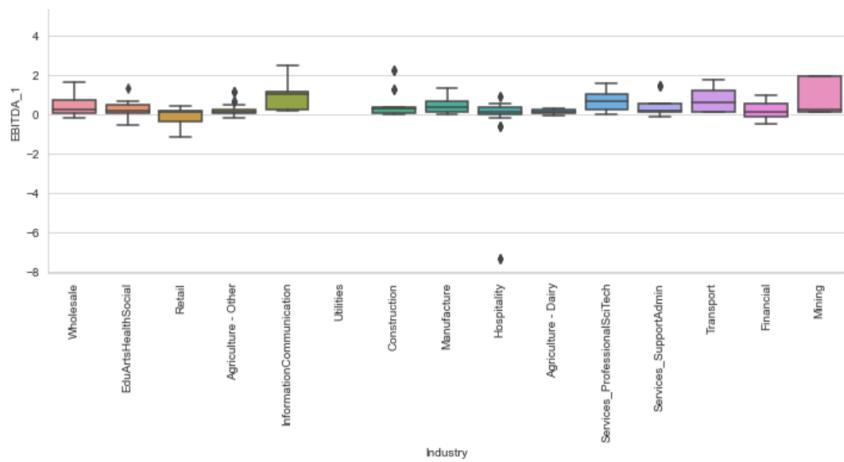
```
[9]: data['Industry'].unique()
```

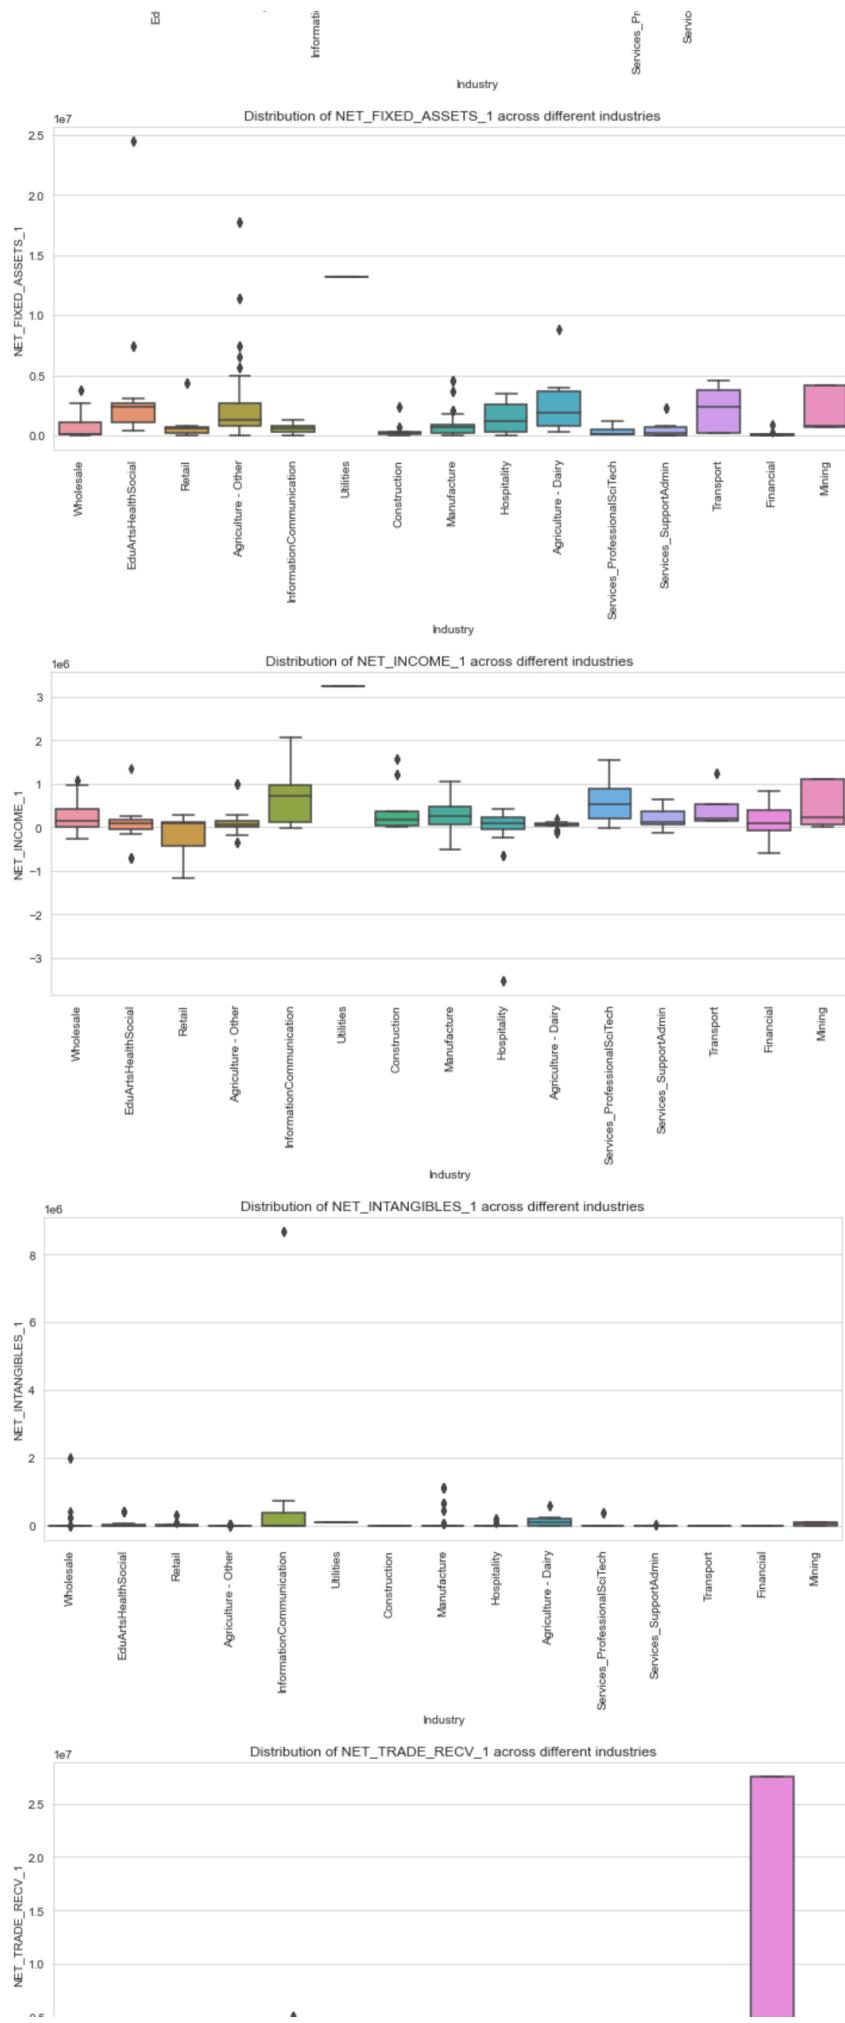
```
[9]: array(['Wholesale', 'EduArtsHealthSocial', 'Retail',
       'Agriculture - Other', 'InformationCommunication', 'Utilities',
       'Construction', 'Manufacture', 'Hospitality',
       'Agriculture - Dairy', 'Services_ProfessionalSciTech',
       'Services_SupportAdmin', 'Transport', 'Financial', 'Mining'],
      dtype=object)
```

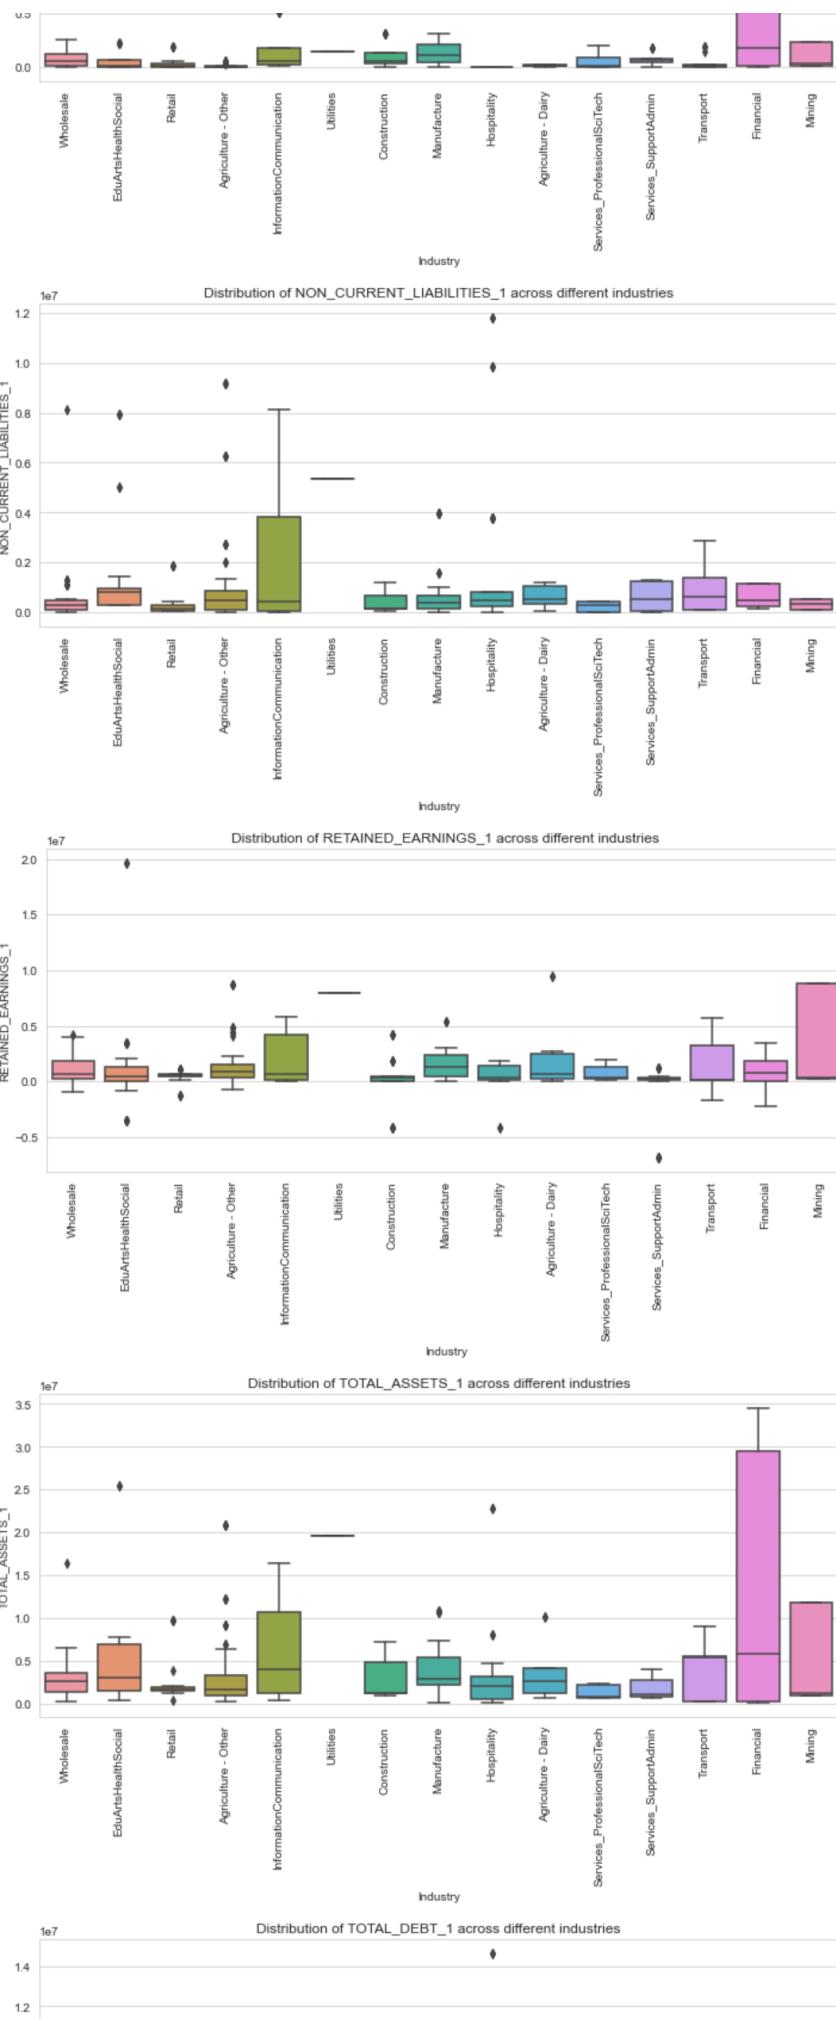
```
[10]: numerical_features = data.select_dtypes(['float64']).columns
for feature in numerical_features:
    plt.figure(figsize=(12,5))
    sns.boxplot(y=data[feature],x=data['Industry'])
    plt.xticks(rotation=90)
    plt.title(f'Distribution of {feature} across different industries')
```

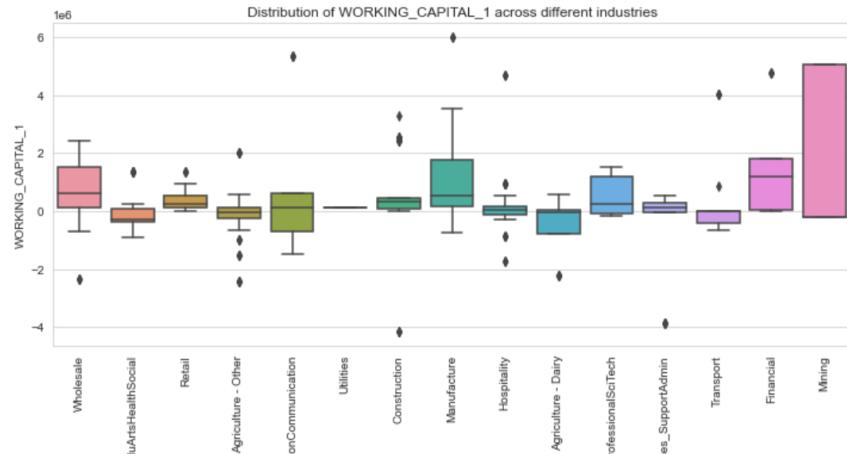
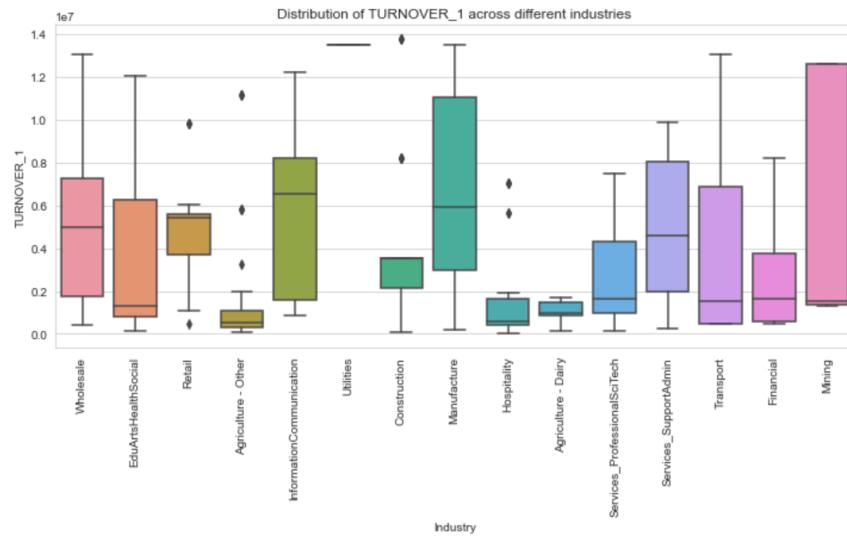
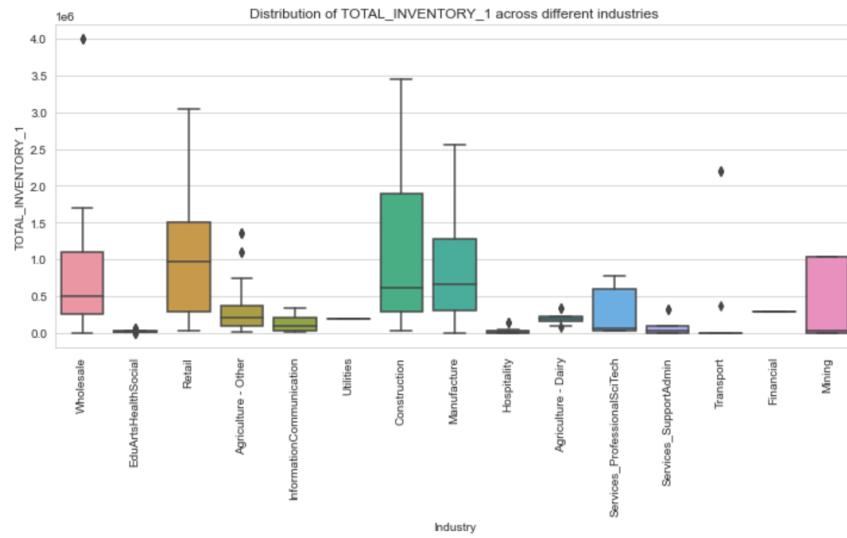
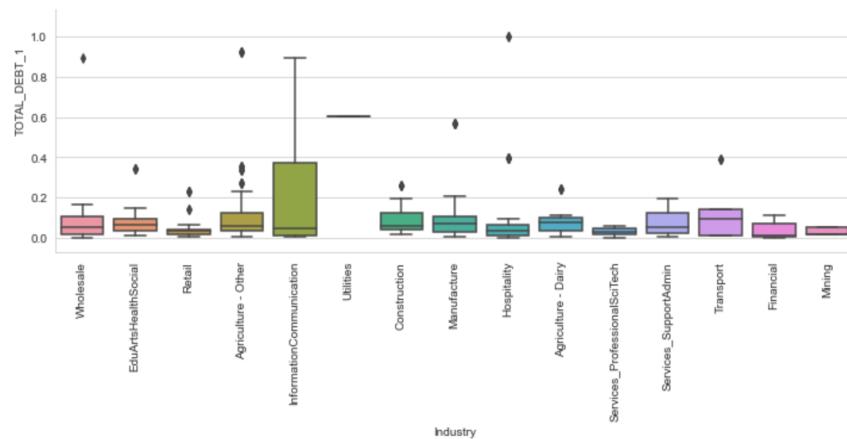


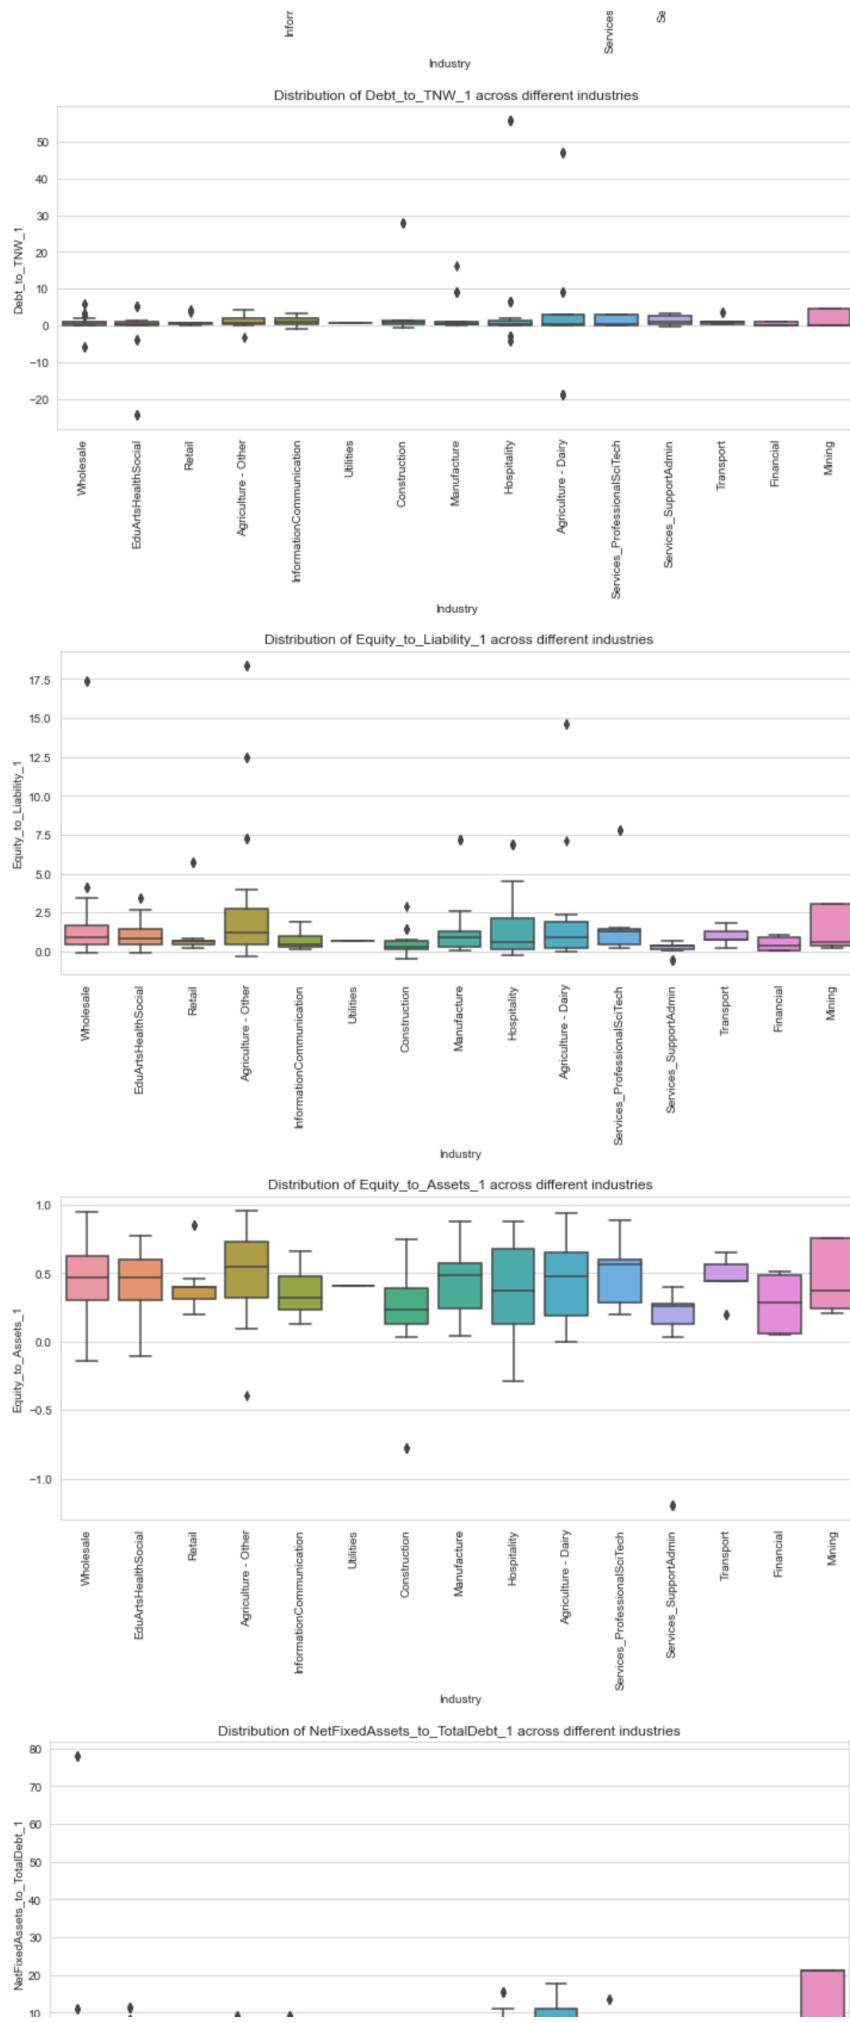


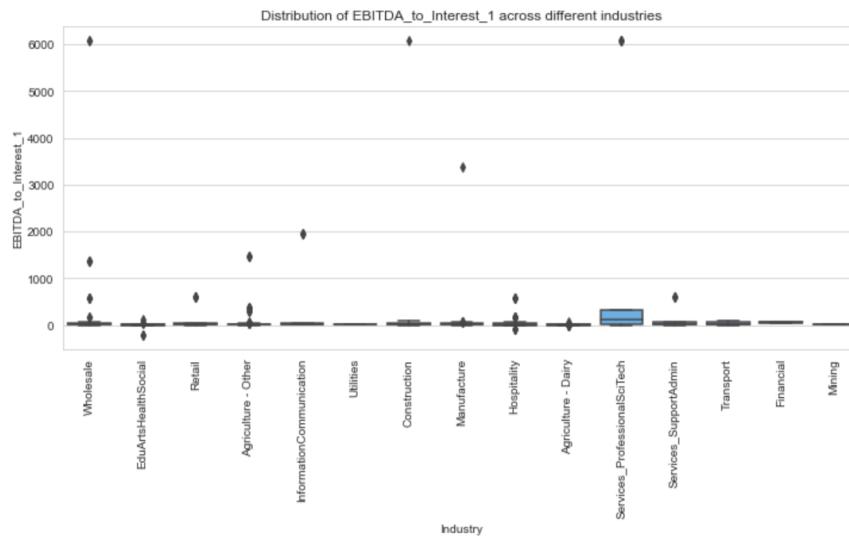
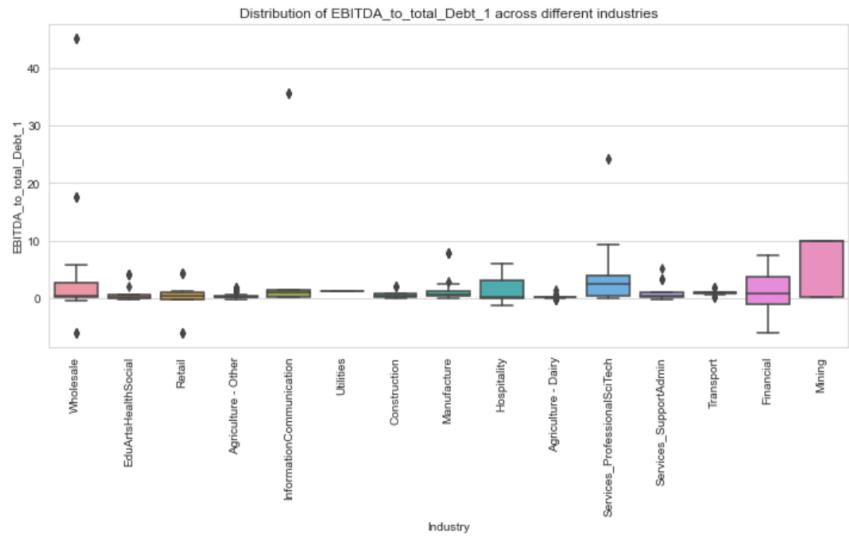
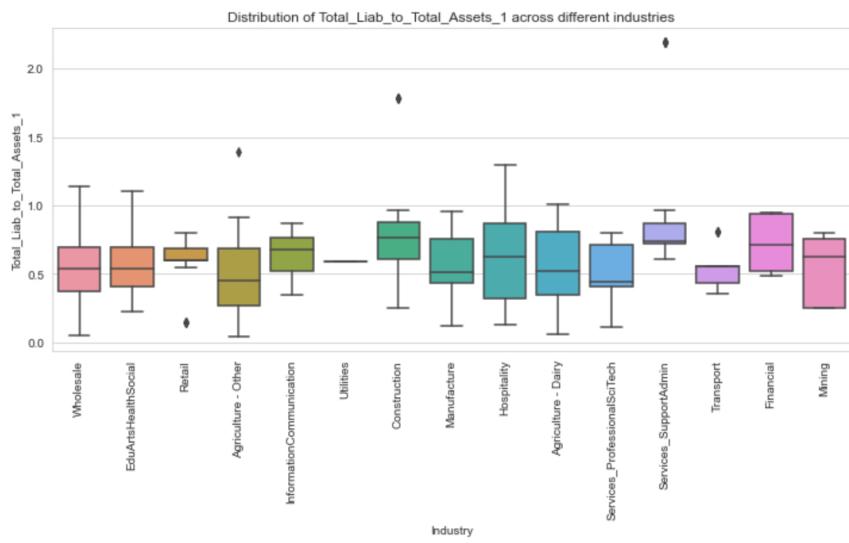
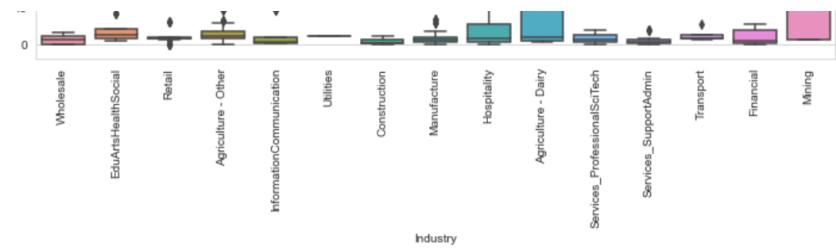


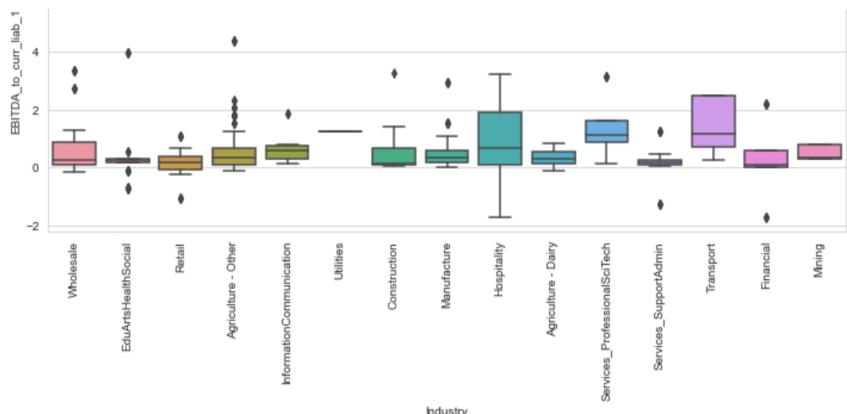




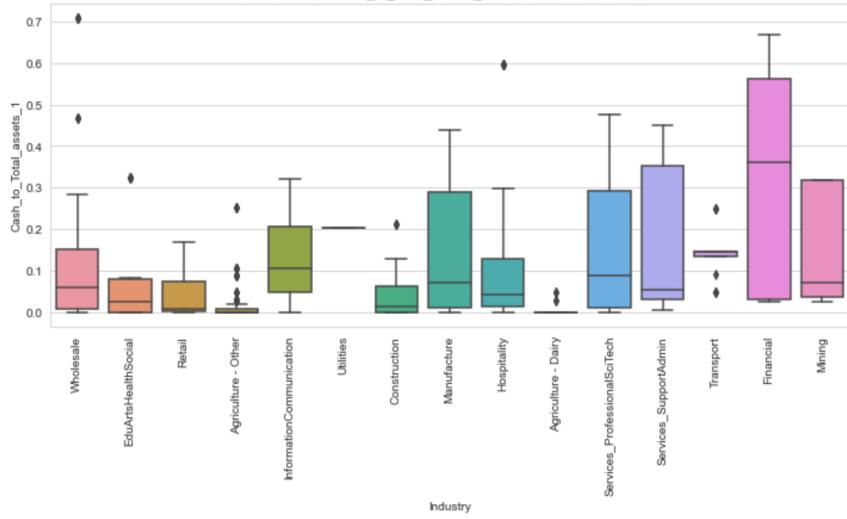




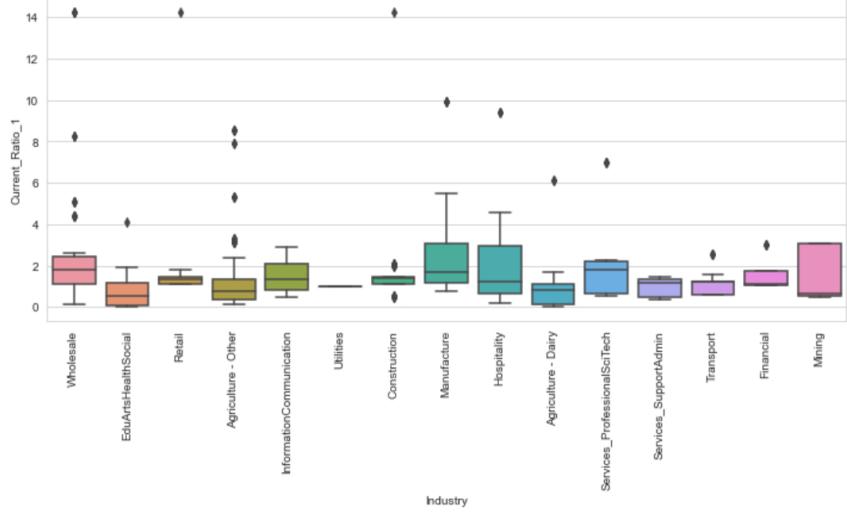




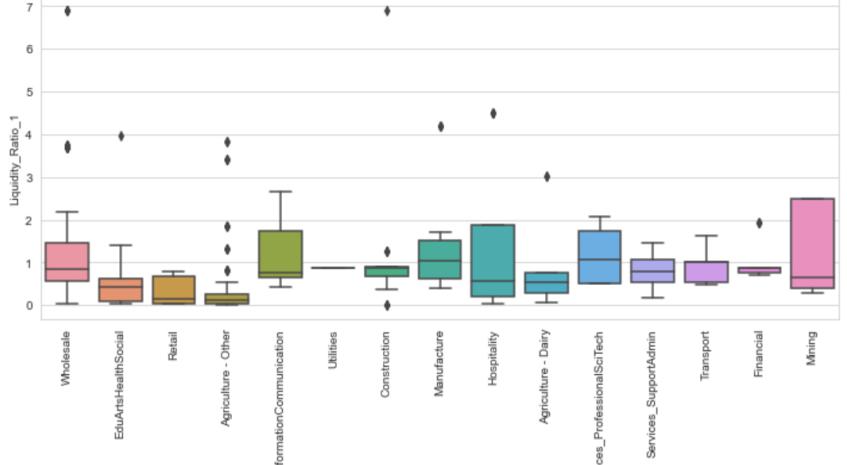
Distribution of Cash\_to\_Total\_assets\_1 across different industries

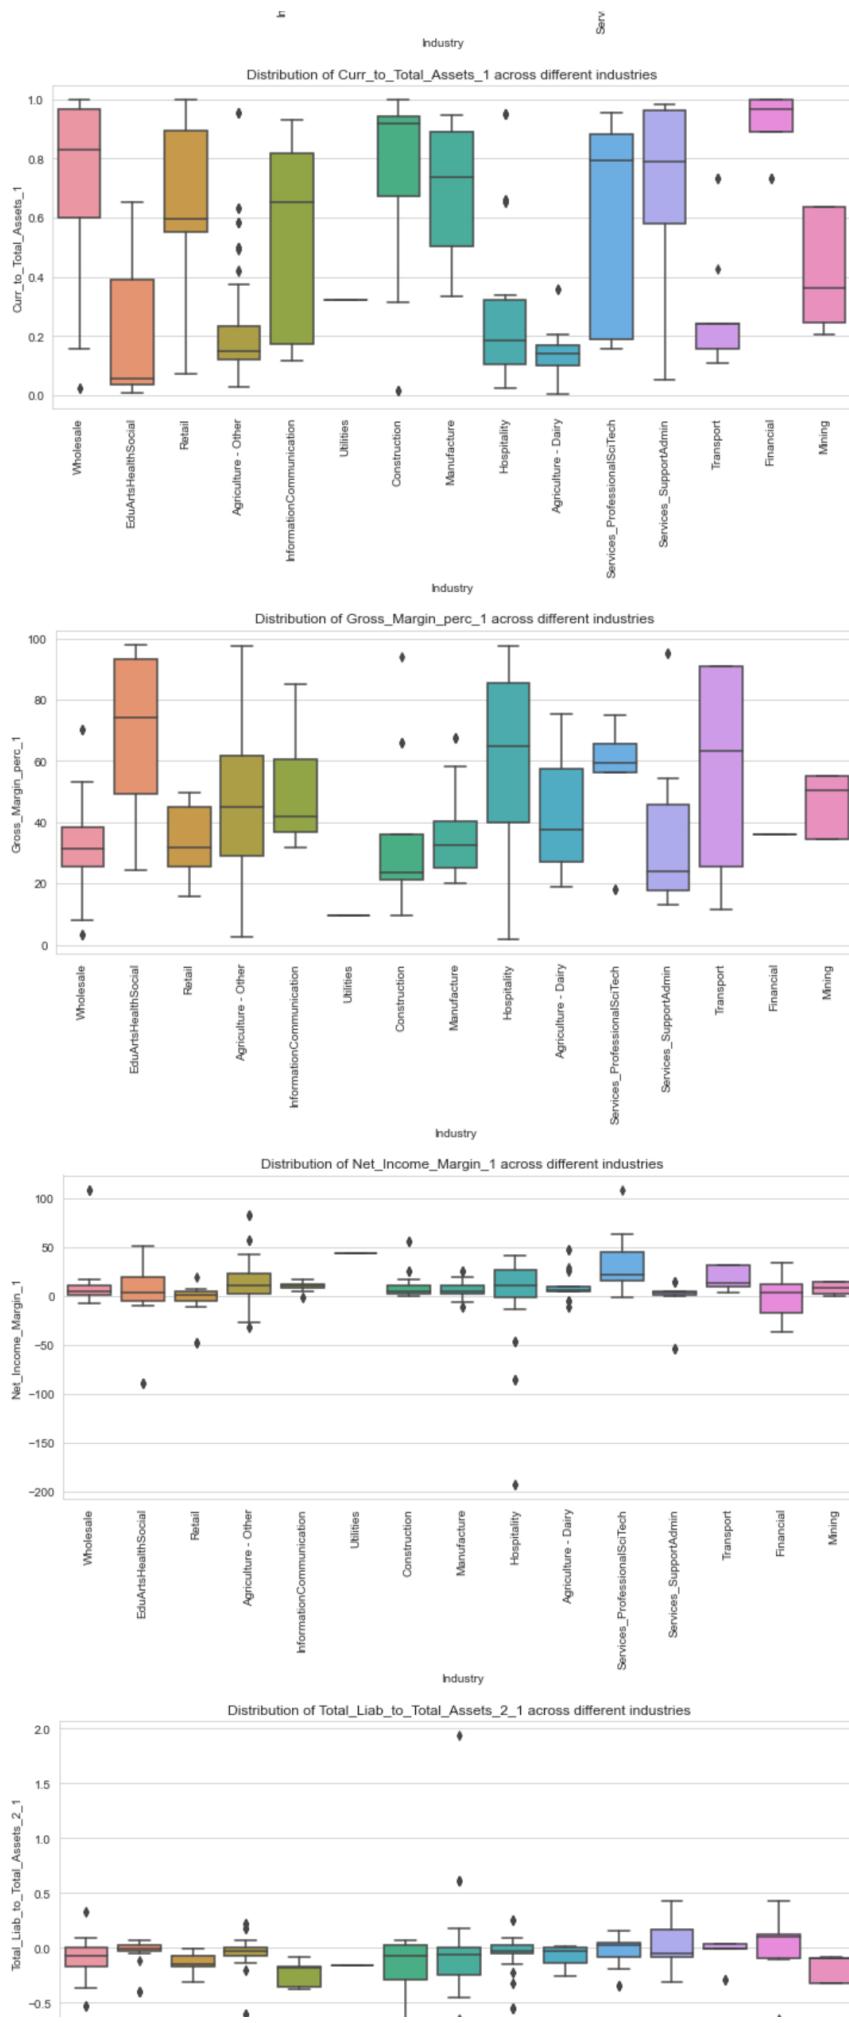


Distribution of Current\_Ratio\_1 across different industries



Distribution of Liquidity\_Ratio\_1 across different industries







- The numerical features varies across different industries
- so we will replace the missing values with the median of their respective industries

```
[11]: #missing value imputation with median
for feature in numerical_features:
    data[feature].fillna(data.groupby('Industry')[feature].transform('median'),inplace=True)

[12]: data.isna().any().any()

[12]: False
```

### Outlier Handling

handling the outliers by quantile based flooring and capping with their respective industries

```
[13]: #outlier treatment
data = data.groupby(data.Industry.values).transform(lambda x: np.clip(x,x.quantile(0.05),x.quantile(0.95))if x.name in numerical_features else x)
data.head()
```

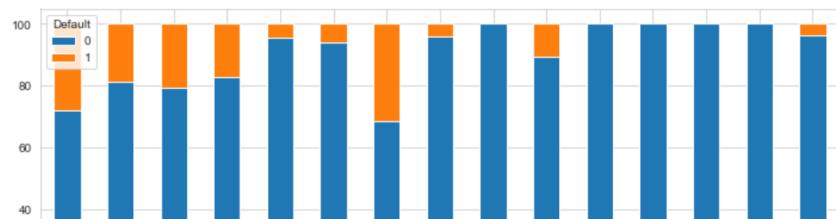
	Industry	CASH_AND_CASH_EQUIVALENTS_floored1	COST_OF_GOODS SOLD_1	CURRENT_ASSETS_1	CURRENT_LIABILITIES_1	DEPRECIATION_1	EBITDA_1	FIXED
0	Wholesale	323773.0	2995379.0	476381.0	516192.0	78247.0	357842.0	
1	EduArtsHealthSocial	18593.0	11753.0	21913.0	257405.0	6988.0	42274.0	
2	Retail	202720.0	3420276.0	665190.0	373352.0	76899.0	406223.0	
3	Agriculture - Other	87515.0	2349333.7	1753108.7	2153519.0	134105.0	657162.4	
4	EduArtsHealthSocial	272938.0	237933.0	381958.0	681622.0	81766.0	-515748.0	

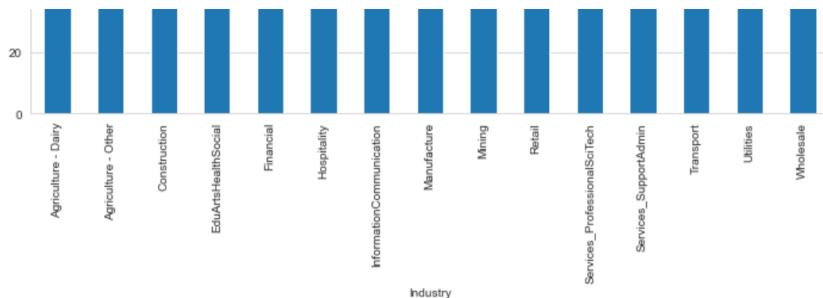
### EDA

```
[15]: prop_by_industry = pd.crosstab(data['Industry'],data['Default']).apply(lambda x:x/x.sum()*100,axis=1)
print(prop_by_industry)
prop_by_industry.plot(kind='bar',stacked=True)
plt.xticks(rotation=90)
```

Default	0	1
Industry		
Agriculture - Dairy	72.000000	28.000000
Agriculture - Other	81.300813	18.699187
Construction	79.310345	20.689655
EduArtsHealthSocial	82.926829	17.073171
Financial	95.454545	4.545455
Hospitality	94.000000	6.000000
InformationCommunication	68.421053	31.578947
Manufacture	95.774648	4.225352
Mining	100.000000	0.000000
Retail	89.285714	10.714286
Services_ProfessionalSciTech	100.000000	0.000000
Services_SupportAdmin	100.000000	0.000000
Transport	100.000000	0.000000
Utilities	100.000000	0.000000
Wholesale	96.153846	3.846154

```
[15]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14]),
 [Text(0, 0, 'Agriculture - Dairy'),
  Text(1, 0, 'Agriculture - Other'),
  Text(2, 0, 'Construction'),
  Text(3, 0, 'EduArtsHealthSocial'),
  Text(4, 0, 'Financial'),
  Text(5, 0, 'Hospitality'),
  Text(6, 0, 'InformationCommunication'),
  Text(7, 0, 'Manufacture'),
  Text(8, 0, 'Mining'),
  Text(9, 0, 'Retail'),
  Text(10, 0, 'Services_ProfessionalSciTech'),
  Text(11, 0, 'Services_SupportAdmin'),
  Text(12, 0, 'Transport'),
  Text(13, 0, 'Utilities'),
  Text(14, 0, 'Wholesale')])
```

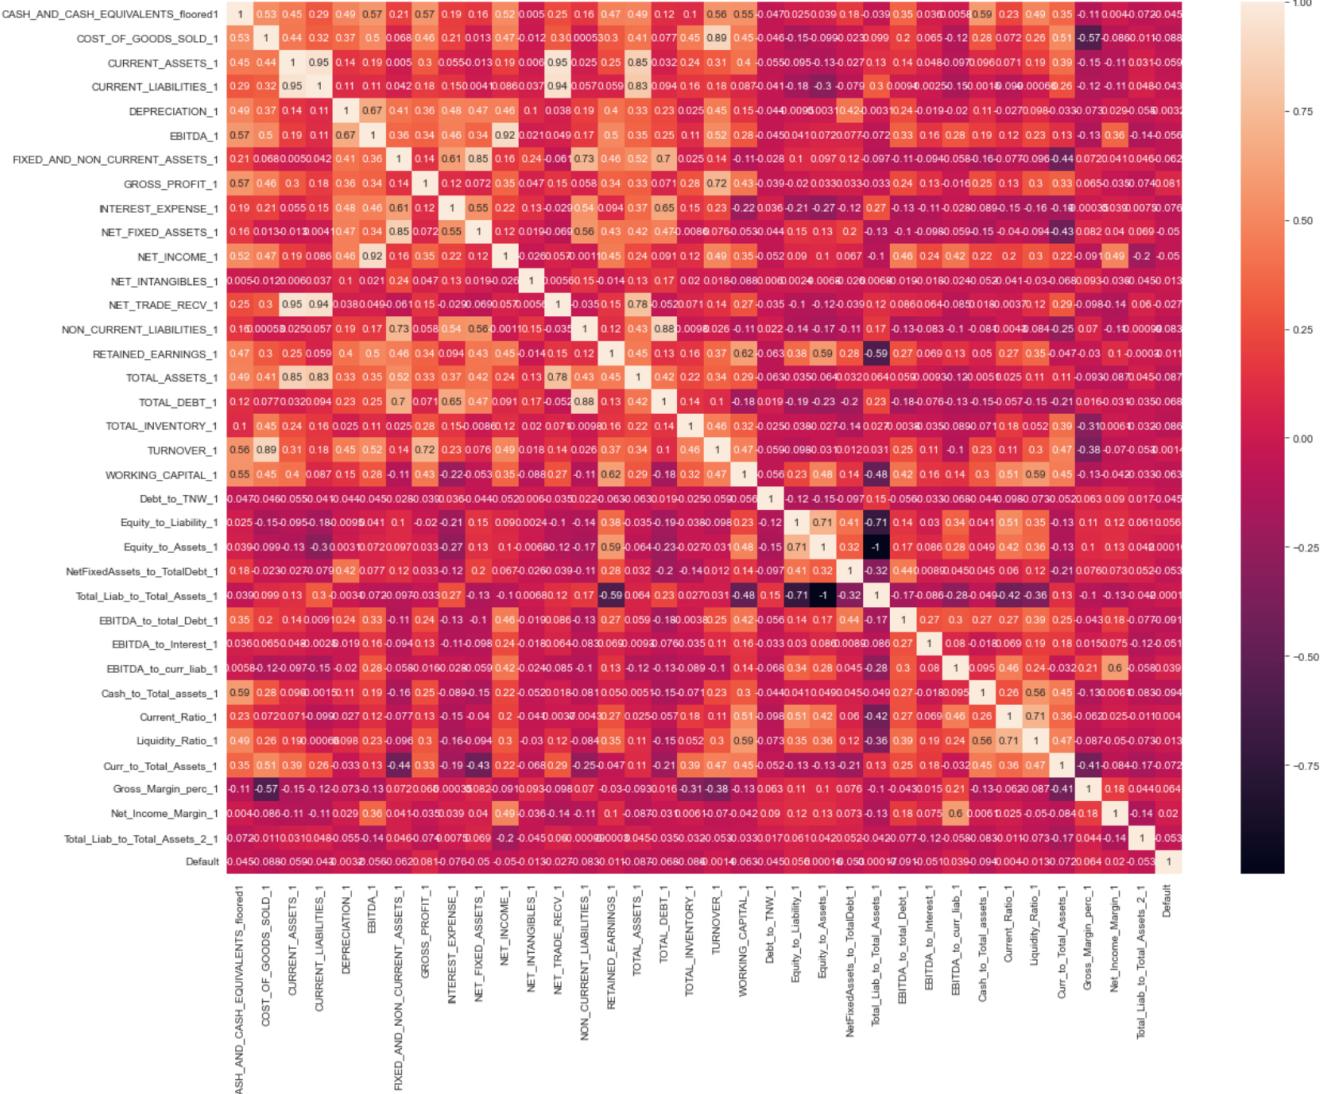




- Among all the industries InformationCommunication and Agriculture-diary has the highest default rate
- Agriculture other , construction and EduArtsHealthSocial have default rate around 20%

```
[16]: plt.figure(figsize=(20,15))
sns.heatmap(data.corr(), annot=True)
```

```
[16]: <AxesSubplot:>
```



some of the features are highly correlated to each other i,e Multicollinearity

```
[17]: #scaling the data using standard scaler

scaler = StandardScaler()
data_trf = data.copy()
data_trf[numerical_features] = scaler.fit_transform(data_trf[numerical_features])
data_trf.head()
```

	Industry	CASH_AND_CASH_EQUIVALENTS_floored1	COST_OF_GOODS SOLD_1	CURRENT_ASSETS_1	CURRENT_LIABILITIES_1	DEPRECIATION_1	EBITDA_1	FIXED
0	Wholesale	0.082289	0.227524	-0.353960	-0.277880	-0.099619	-0.017990	
1	EduArtsHealthSocial	-0.293261	-0.877496	-0.472153	-0.350854	-0.563862	-0.460186	
2	Retail	-0.066677	0.384889	-0.304856	-0.318159	-0.108401	0.049805	
3	Agriculture - Other	-0.208447	-0.011747	-0.021921	0.183818	0.264288	0.401438	

```
4 EduArtsHealthSocial          0.019732         -0.793728        -0.378516        -0.231232        -0.076693   -1.242125
```

```
[18]: #seperating independent and dependent variables
X = data_trf.loc[:,data_trf.columns != 'Default']
X.drop(['Industry'],axis=1,inplace=True)
y = data_trf['Default']
print('shape of X',X.shape)
print('shape of y',y.shape)

shape of X (573, 35)
shape of y (573,)

[19]: # using variance inflation factor vif for removing highly correlated features
# using a threshold of 4 for VIF
from statsmodels.stats.outliers_influence import variance_inflation_factor
features = list(X.columns)
while True:
    vif = pd.DataFrame()
    df = X[features]
    vif['variables'] = df.columns
    vif['VIF'] = [variance_inflation_factor(df.values,i)for i in range(df.shape[1])]
    vif.sort_values(by='VIF',ascending=False,inplace=True,ignore_index=True)
    if vif['VIF'][0] > 4:
        print(vif['variables'][0], 'has been removed with a vif of' , vif['VIF'][0])
        features.remove(vif['variables'][0])
    else:
        break
print('NO of Remaining features after VIF',len(features))

Equity_to_Assets_1 has been removed with a vif of 14735685.088588918
CURRENT_LIABILITIES_1 has been removed with a vif of 759.1328261534444
CURRENT_ASSETS_1 has been removed with a vif of 247.62206643656504
TOTAL_ASSETS_1 has been removed with a vif of 73.88002943534978
EBITDA_1 has been removed with a vif of 32.962967868964235
TURNOVER_1 has been removed with a vif of 19.8759711335175
FIXED_AND_NON_CURRENT_ASSETS_1 has been removed with a vif of 13.991232694600122
TOTAL_DEBT_1 has been removed with a vif of 8.815008158029718
WORKING_CAPITAL_1 has been removed with a vif of 5.551574543389508
Current_Ratio_1 has been removed with a vif of 4.611448736375268
NET_INCOME_1 has been removed with a vif of 4.4843431164759755
CASH_AND_CASH_EQUIVALENTS_floor01 has been removed with a vif of 4.047231973943805
NO of Remaining features after VIF 23
```

```
[20]: X =X[features]
X.shape
```

```
[20]: (573, 23)
```

## MODEL BUILDING

```
[21]: #building the baseline model

x_train,x_test,y_train,y_test=train_test_split(X,y,test_size=0.1,random_state=122)
print("Target Class distribution in Training split:",Counter(y_train))
print("Target Class distribution in Testing split:",Counter(y_test))

Target Class distribution in Training split: Counter({0: 460, 1: 55})
Target Class distribution in Testing split: Counter({0: 51, 1: 7})
```

```
[22]: def metrics(model,testdata,actual,predicted):
    ''' This function returns evaluation metrics by comparing actual and predicted outcomes'''
    print('Testing F1-Score:',f1_score(actual,predicted))
    probpred = model.predict_proba(testdata)
    print('AUC Score:',roc_auc_score(actual,probpred[:,1]))
    print('Classification Report \n',classification_report(actual,predicted))
    fpr , tpr ,thresh = roc_curve(actual,probpred[:,1],pos_label=1)
    plt.figure(figsize=(12,5))
    plt.subplot(1,2,1)
    sns.heatmap(confusion_matrix(actual,predicted),annot=True,fmt=' .1f')
    plt.title('Confusion matrix')
    plt.subplot(1,2,2)
    plt.plot(fpr,tpr,linestyle='--',color='r')
    plt.title('ROC Curve')
```

## Training the model

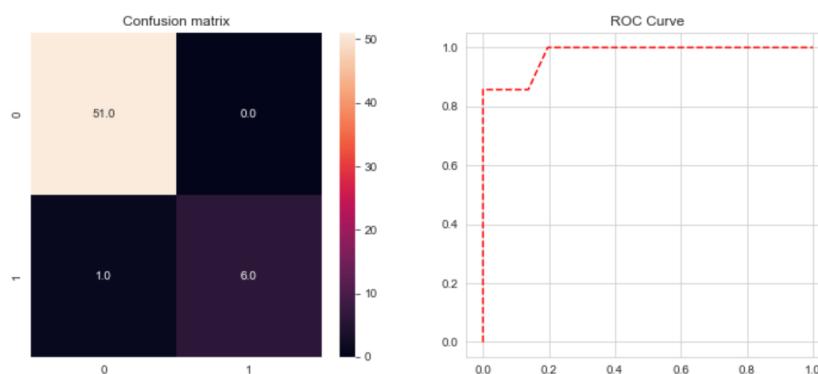
```
[23]: classifier = RandomForestClassifier()
classifier.fit(x_train,y_train)
Trainpred = classifier.predict(x_train)
print('Training F1-Score:',f1_score(y_train,Trainpred))

Training F1-Score: 0.9734513274336283
```

```
[24]: #predicting on the test data
pred = classifier.predict(x_test)
metrics(classifier,x_test,y_test,pred)

Testing F1-Score: 0.923076923076923
AUC Score: 0.9761984761904762
Classification Report
      precision    recall   f1-score   support
       0         0.98     1.00     0.99      51
       1         1.00     0.86     0.92       7
```

accuracy		0.98	58
macro avg	0.99	0.93	0.96
weighted avg	0.98	0.98	0.98



- The testing F1-score is 92% which is good but,
- the training F1-Score is higher i.e 98%, this shows that our model is slightly overfitting

### K-FOLD CROSS VALIDATION to reduce the overfitting

```
[25]: cv = StratifiedKFold(n_splits=5,random_state=22,shuffle=True)
scores = cross_val_score(classifier,X,y,scoring='f1',n_jobs=-1)
print("F1_Score:",np.mean(scores),"std:",np.std(scores))

F1_Score: 0.9426666666666668 std: 0.05681940200710008
```

- The testing f1-score now has been improved from 92% to 94.1%
- This slightly reduces the problem of overfitting

### Hyperparameter tuning

```
[26]: params=param_grid = {'bootstrap': [True],
'max_depth': [1,4,10, 20, 30, 40, 50, 60, 70, 80, 90, 100, None],
'max_features': ['auto', 'sqrt'],
'min_samples_leaf': [1, 2, 4],
'min_samples_split': [2, 5, 10],
'n_estimators': [200, 400, 600, 800, 1000]
}
grid = GridSearchCV(classifier,params,scoring='f1',n_jobs=-1,cv=5,verbose=2)
grid.fit(X,y)

Fitting 5 folds for each of 1170 candidates, totalling 5850 fits
[26]: GridSearchCV(cv=5, estimator=RandomForestClassifier(), n_jobs=-1,
      param_grid=[{'bootstrap': [True],
                   'max_depth': [1, 4, 10, 20, 30, 40, 50, 60, 70, 80, 90,
                   100, None],
                   'max_features': ['auto', 'sqrt'],
                   'min_samples_leaf': [1, 2, 4],
                   'min_samples_split': [2, 5, 10],
                   'n_estimators': [200, 400, 600, 800, 1000]},
                  scoring='f1', verbose=2)

[27]: print('Best params:',grid.best_params_)
print('Best Score:',grid.best_score_)
grid.best_estimator_

Best params: {'bootstrap': True, 'max_depth': 20, 'max_features': 'auto', 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 400}
Best Score: 0.9426666666666668
[27]: RandomForestClassifier(max_depth=20, n_estimators=400)
```

```
[28]: train_pred = grid.predict(x_train)
test_pred = grid.predict(x_test)
print('Training score after hyperparameter tuning:',f1_score(y_train,train_pred))
metrics(grid,x_test,y_test,test_pred)
```

Training score after hyperparameter tuning: 0.9734513274336283

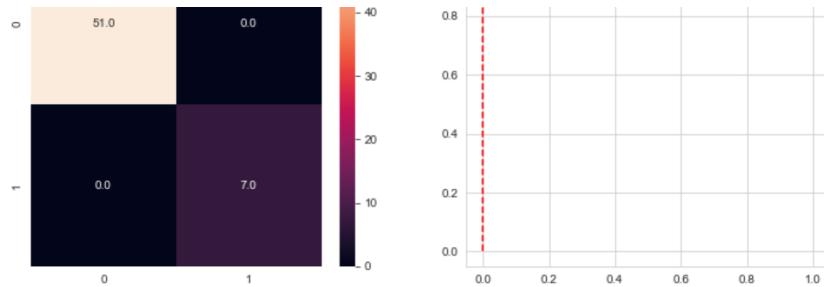
Testing F1-Score: 1.0

AUC Score: 1.0

Classification Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	51
1	1.00	1.00	1.00	7
accuracy			1.00	58
macro avg	1.00	1.00	1.00	58
weighted avg	1.00	1.00	1.00	58

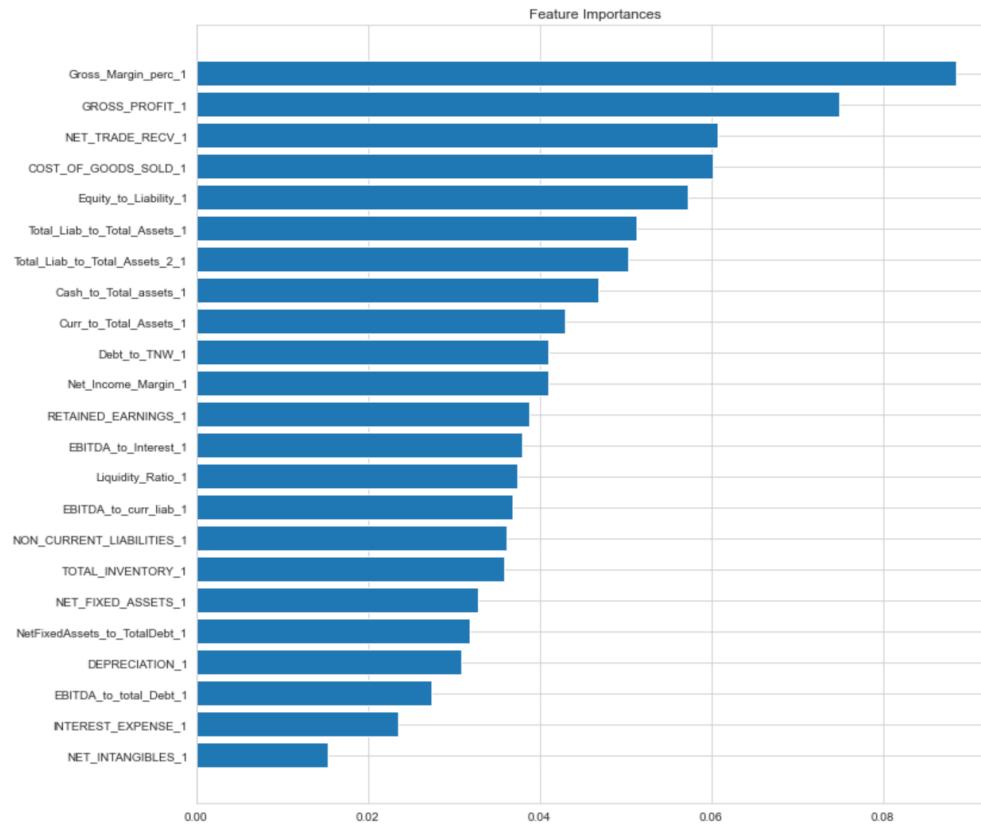




- The testing f1\_score improved gradually from 94% to 99%
- The Training f1\_score is 97%
- The overfitting problem has been removed

```
[29]: feat_imps=pd.DataFrame(classifier.feature_importances_,index=x_train.columns,columns=['Scores'])
feat_imps.sort_values(by='Scores',inplace=True,ascending=True)
plt.figure(figsize=(12,12))
plt.barh(width=feat_imps['Scores'],y=feat_imps.index)
plt.title('Feature Importances')
```

[29]: Text(0.5, 1.0, 'Feature Importances')



```
[32]: print('The most important features for predicting the default are:',list(feat_imps.index[:5]))
```

The most important features for predicting the default are: ['NET\_INTANGIBLES\_1', 'INTEREST\_EXPENSE\_1', 'EBITDA\_to\_total\_Debt\_1', 'DEPRECIATION\_1', 'NetFixedAssets\_to\_TotalDebt\_1']